

EDUCATION

Princeton University

M.S.E, Computer Science |

Princeton, NJ

August 2024-Present

Funded by the machine learning department under the supervision of Prof. Tom Griffiths in the computational cognitive science [lab](#).

Relevant Coursework: Machine Learning and Pattern Recognition, Foundations of Probabilistic Modeling, AI Safety, Cognitive Psychology

Research Interests: Relational reasoning, metalearning, and human-like learning in neural networks

Columbia University

B.A, Applied Mathematics | GPA: 3.51 | Major GPA: 3.71

New York, NY

September 2017-May 2021

Relevant Coursework: Natural Language Processing, Applied Deep Learning, Data Structures, Discrete Math, Linear Algebra, Intro to Statistics, Discrete Math, Complex Variables, Analysis and Optimization, Calculus III & IV

PRESENTATIONS & PAPERS

- **Gupta M.**, Rane S., McCoy T., Griffiths T. *Convolutional Neural Networks Can (Meta)-Learn the Same-Different Relation*: In the Proceedings of the 47th Annual Conference of the Cognitive Science Society
- Bencomo G., **Gupta M.**, Marinescu I., McCoy T., Griffiths T. *Teasing Apart Architecture and Initial Weights as Sources of Inductive Bias in Neural Networks*: In the Proceedings of the 47th Annual Conference of the Cognitive Science Society
- **Gupta M.** *Navigating High Dimensional Concept Space with Metalearning*: ICML 2025, Workshop on High-Dimensional Learning Dynamics
 - **Gupta, M.**, Franke M., Hawkins R., Wu C. *Pragmatic Vigilance Inoculates Social Networks Against Misinformation*, Presentation to the Computational Pragmatics Lab: University of Tübingen
- Azizi E., Azad T., **Gupta M.**, Nazaret A. *Ensembling in Variational Autoencoder Architectures for Effective Posterior Distribution of Cell State Estimation*. Pre-print and Azizi Lab Presentation: Columbia University
- **Gupta M.**, Nika J., Carsault T. [Multi-Step Chord Prediction for Human-AI Improvisation](#). Presentation: Columbia Undergraduate Research Symposium, New York, NY (Virtual): September 2020
- **Gupta M.**, Malt M. [Musical Markov Chains](#). Presentation: Reid Hall Research Seminar, Paris, France: May 2020

AWARDS & HONORS

- Hackathon Winner: Cognitive Modeling of Humans vs. Large Language Models, University of Tübingen – 2023
 - Dean's List – Columbia University 2020-2021
 - Heinrich Research Fellowship, Columbia University (\$2500) - 2020
 - Spritz Family Research Grant, Columbia University (\$3000) – 2020
 - PSAT National Commended Scholar – 2017
 - Canadian Gold Medalist in both French (2016-2017) and Latin (2014-2017) National Language Exams
 - Greville Smith Scholarship, McGill University (\$48,000, not availed) – 2017
- Hugh M. Brock National Entrance Scholarship, University of British Columbia (\$30,400, not availed) - 2017

RESEARCH EXPERIENCE

Max Planck Institute – Center For Humans and Machines

Berlin, Germany

AI Behaviorist - Research Assistant

May-August 2024

- Engineered online JS experiments to simulate multi-generational human-Large Language Model (LLM) communications in transmission networks of varying sizes and types. Built for research published in human-AI interaction conferences.
- Contributed visualizations and NLP analyses of public sentiments on AI risks/rewards across 400 human participants from India and the US. Compared attitudes via sociological coding.

University of Tübingen

Tübingen, Germany

Research Intern – Supervisors: Michael Franke, Robert Hawkins, Charley Wu

Oct 2023 - Present

- Built a multi-agent extension of the Rational Speech Act (RSA) framework to examine effects of persuasive social goals

on belief diffusion in social networks. Simulating RSA communications in a random network in Python and R and examining belief convergence with Bayesian inference over iterated dyadic communications.

- First-authored a paper illustrating conditions under which listener vigilance inoculates networks against misinformation.

Columbia University Medical Center

New York, NY

Research Assistant - Supervisor: Elham Azizi

Jan – August 2021

- Analyzed the effects of aggregating neural network outputs to form posterior distributions (ensembling). Compared deep ensembling and batch ensembling on variational autoencoders (VAE's) performing differential gene expression.
- Trained Bayesian neural networks and VAE's on single-cell data in PyTorch analyzed posteriors across random initializations. Presented methods and findings to the computational cancer biology lab.

IRCAM, Centre Pompidou

Paris, France

Research Fellow – Supervisors: Mikhail Malt, Jérôme Nika

Jan-Sep 2020

- Built language models (RNNs, LSTMs) trained in PyTorch on musical data from live jazz for human-AI improvisations.
- Authored, presented a report and poster on generative music with language models at Columbia's 2020 research symposium. Authored a manuscript on stochastic models in music, published at Reid Hall's research symposium.

PROFESSIONAL EXPERIENCE

Weill Cornell Medical, Cornell University

New York, NY

Research Software Engineer

May 2022 - Nov 2022

- Worked on efficiency and indexing of the main NLP pipeline, using OCR to parse doctor notes into machine-readable text.
- Built an AWS-hosted ETL pipeline with Docker, Python, Java, and SQL to securely geocode address data from hospital patients, increasing geocoding accuracy and runtime efficiency by 15%.

Infosys Consulting

New York, NY

Business Analyst – AI & Automation

August 2021-May 2022

- Built and deployed an NLG model from open-source to automate 85% of credit loan risk report writing at a top 3 US bank.
- Engineered several NLP models for financial document classification, summarization, and generation in NLTK and JS.
- Assisted executive advising at 2 of the top 5 US banks in AI automation, chatbot implementation, and process mining. Wrote concise technical guides on each of the above topics, shared to clients and the firm at large.

TEACHING EXPERIENCE

Assistant Instructor: COS 126, COS 240

Princeton, NJ

Princeton University

Sep. 2024 – Present

- Weekly teaching and grading for a cohort of 20 Princeton undergraduates with coding and theory p-sets and exams. COS 126 is an intro to java class and COS 240 is a proof-based math class for CS majors.

Teaching Assistant: Reinforcement Learning for Language Model Training

Tübingen, Germany

University of Tübingen

Nov. 2023 – Present

- Grading and coding support for 50+ Tübingen graduate students using Tensorflow for p-sets and RL research projects.

CompTIA

New York, NY

Head Data Science Instructor

Jan 2022-Present

- Head instructor for an online data science and coding boot-camp for Python and SQL. Assistant teacher for web development with React JS: providing comprehensive grading, coding, and career support to students.
- Design and deliver all written and technical curriculum on computer science, statistics, and data analysis, focused on libraries like Pandas, NumPy and Matplotlib. I hand-write the content for the course in Python/Markdown.

Teaching Assistant: Calculus IV

New York, NY / Remote

Columbia University

Jan - May 2021

- Graded assignments and held office hours for ~80 students in Professor Daniela De Silva's class. Coordinated grading across sections. Wrote technical guides and explanations on calculus-related concepts from the course textbooks.

EXTRACURRICULAR EDUCATION

Journal Clubs: Meta-Learning and Mechanistic Interpretability

Princeton, New Jersey

Organized journal clubs for post-doc and graduate students to come together bi-weekly to discuss topics in meta-learning and mechanistic interpretability (with support of the Natural and Artificial Minds initiative at Princeton).

National Deep Inference Fabric (NDIF) Pilot

Virtual, Feb 2025

Running mechanistic interpretability analyses on Llama 405-B with NN-sight, an NDIF-run library for mech-interp.

IICCSSS 2023 (Interdisciplinary Computational Cognitive Science Summer School)

Tübingen, Germany

1st Place: Hackathon for Cognitive Modeling

September 2023

Coursework: Comparing language models to humans; Computational modeling for learning; Human language models

ESSLLI 2023 (European Summer School in Language, Logic & Information)

Ljubljana, Slovenia

Coursework: Probabilistic Language of Thought; Formal Language Theory and Neural Networks;

August 2023

Deep Language Learning from Raw Speech; Logic, Data, Examples, and Learning

Center for AI Safety

Berkeley, CA (Remote)

Intro to Machine Learning Safety Fellow

June – August 2023

Grant-funded student. Coursework covering mechanistic interpretability, machine ethics, systemic AI safety, adversarial robustness, and preventing existential risk from future AI systems.

SKILLS, LANGUAGES & INTERESTS

Programming Languages: Python, Java, R, SQL, JavaScript, HTML, Bash, MATLAB

Technical Skills: Machine learning, data visualization, data analysis and statistical insights, computational modeling, scientific communication, experimental design, web design, teaching and curriculum design

Spoken Languages: English (fluent); French (intermediate) **Interests:** Cognitive science, Running, Tennis, Literature